Set up and Access an Ubuntu Data Science VM and Creating a Jupyter Notebook

Prerequisites

The following are required to complete this hands-on lab:

- An active Microsoft Azure subscription.
- An Xfce remote-desktop client such as X2Go

Exercise 1: Create an Ubuntu Data Science VM

The Ubuntu Data Science Virtual Machine for Linux is a virtual-machine image that makes it easy to get started with data science. Multiple tools are already built, installed, and configured in order to get you up and running quickly. The NVIDIA GPU driver, NVIDIA CUDA, and NVIDIA CUDA Deep Neural Network library (cuDNN) are also included, as are Jupyter and several sample Jupyter notebooks. All installed frameworks are GPU-enabled but work on CPUs as well. In this exercise, you will create an instance of the Data Science Virtual Machine for Linux in Azure.

1. Open the Azure Portal in your browser. If asked to log in, do so using your Microsoft account.

2. Click **+ Create a resource** in the menu on the left side of the portal, and then type "data science" (without quotation marks) into the search box. Select **Data Science Virtual Machine for Linux (Ubuntu)** from the results list.

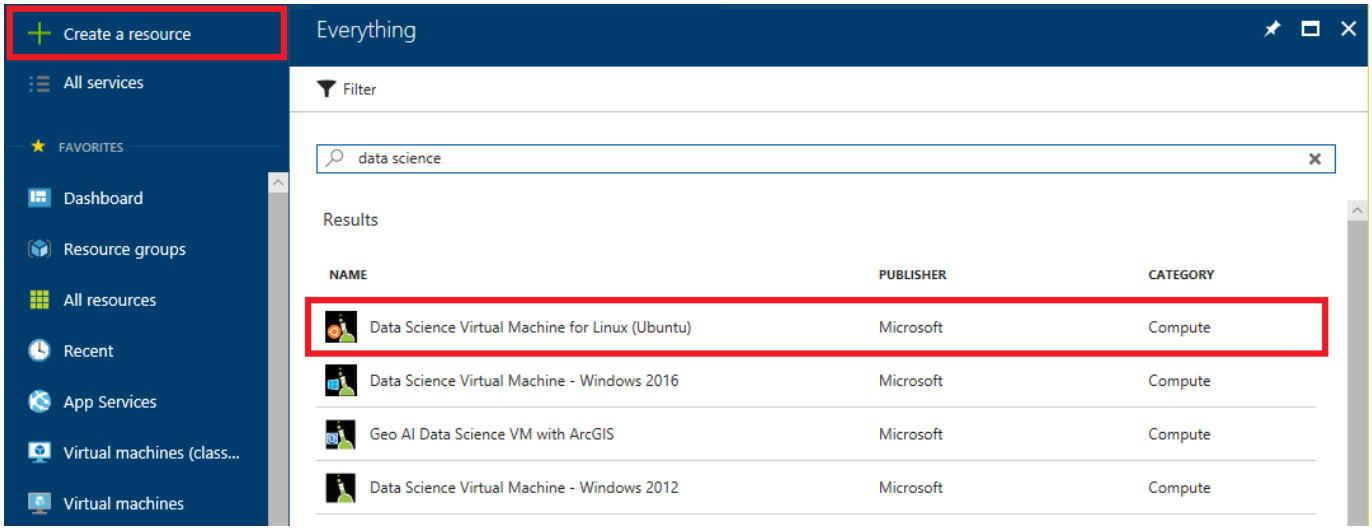

Finding the Ubuntu Data Science VM

3. Take a moment to review the list of tools included in the VM. Then click **Create**.

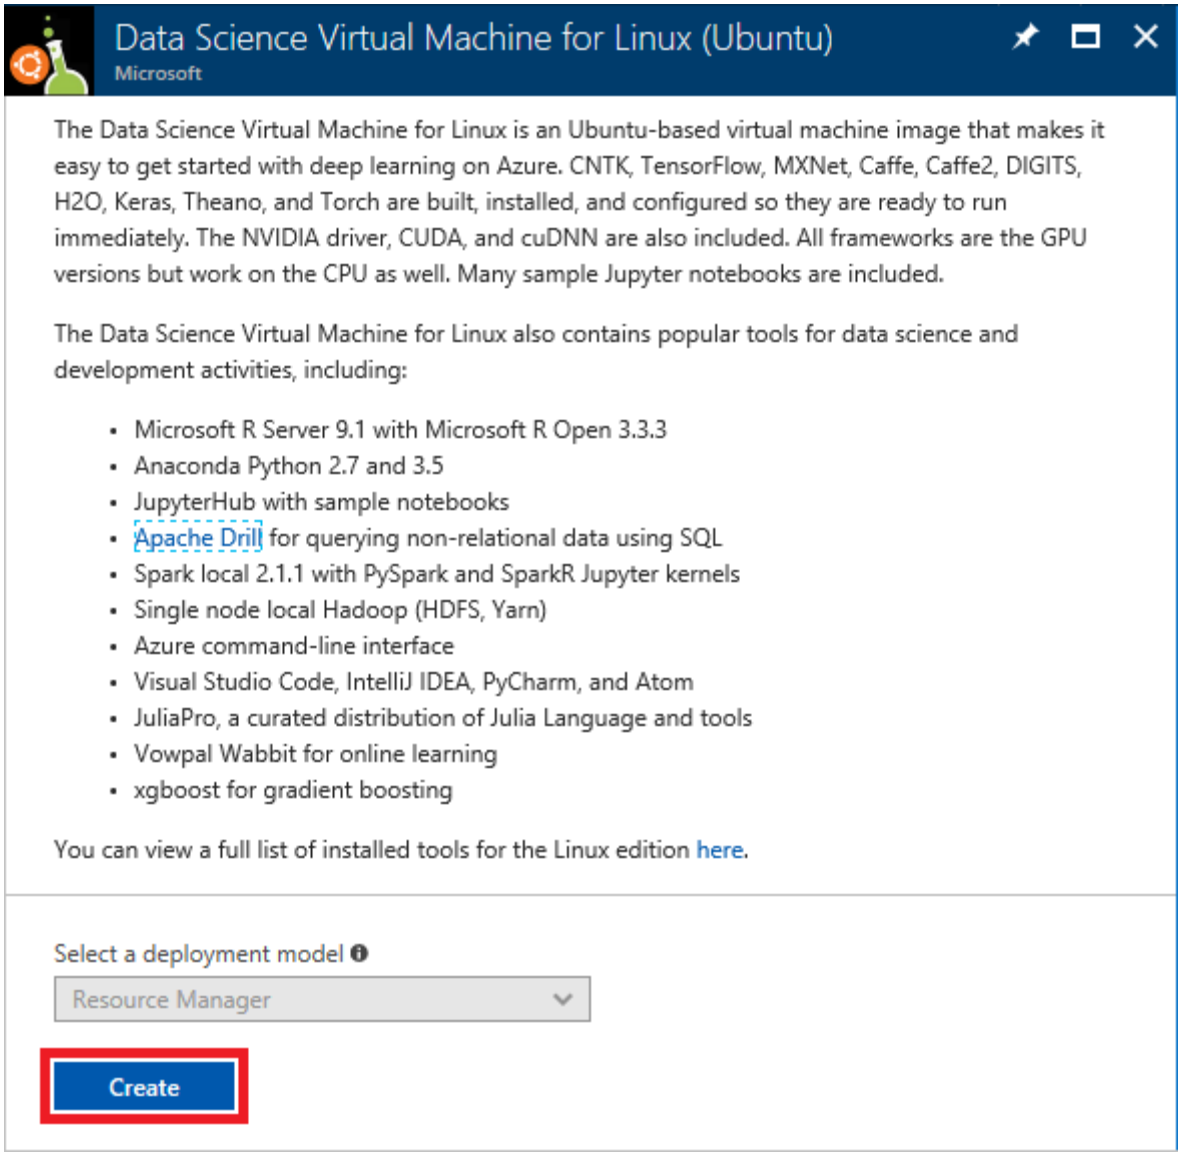

Creating a Data Science VM

4. Enter a name for the virtual machine and a user name for logging into it. Set **Authentication type** to **Password** and enter a password. *Be sure to remember the user name and password that you enter*, because you will need them to access the VM. Select **Create new** under **Resource group** and enter a resource-group name such as "data-science-rg." Select the **Location** nearest you, and then click **OK**..

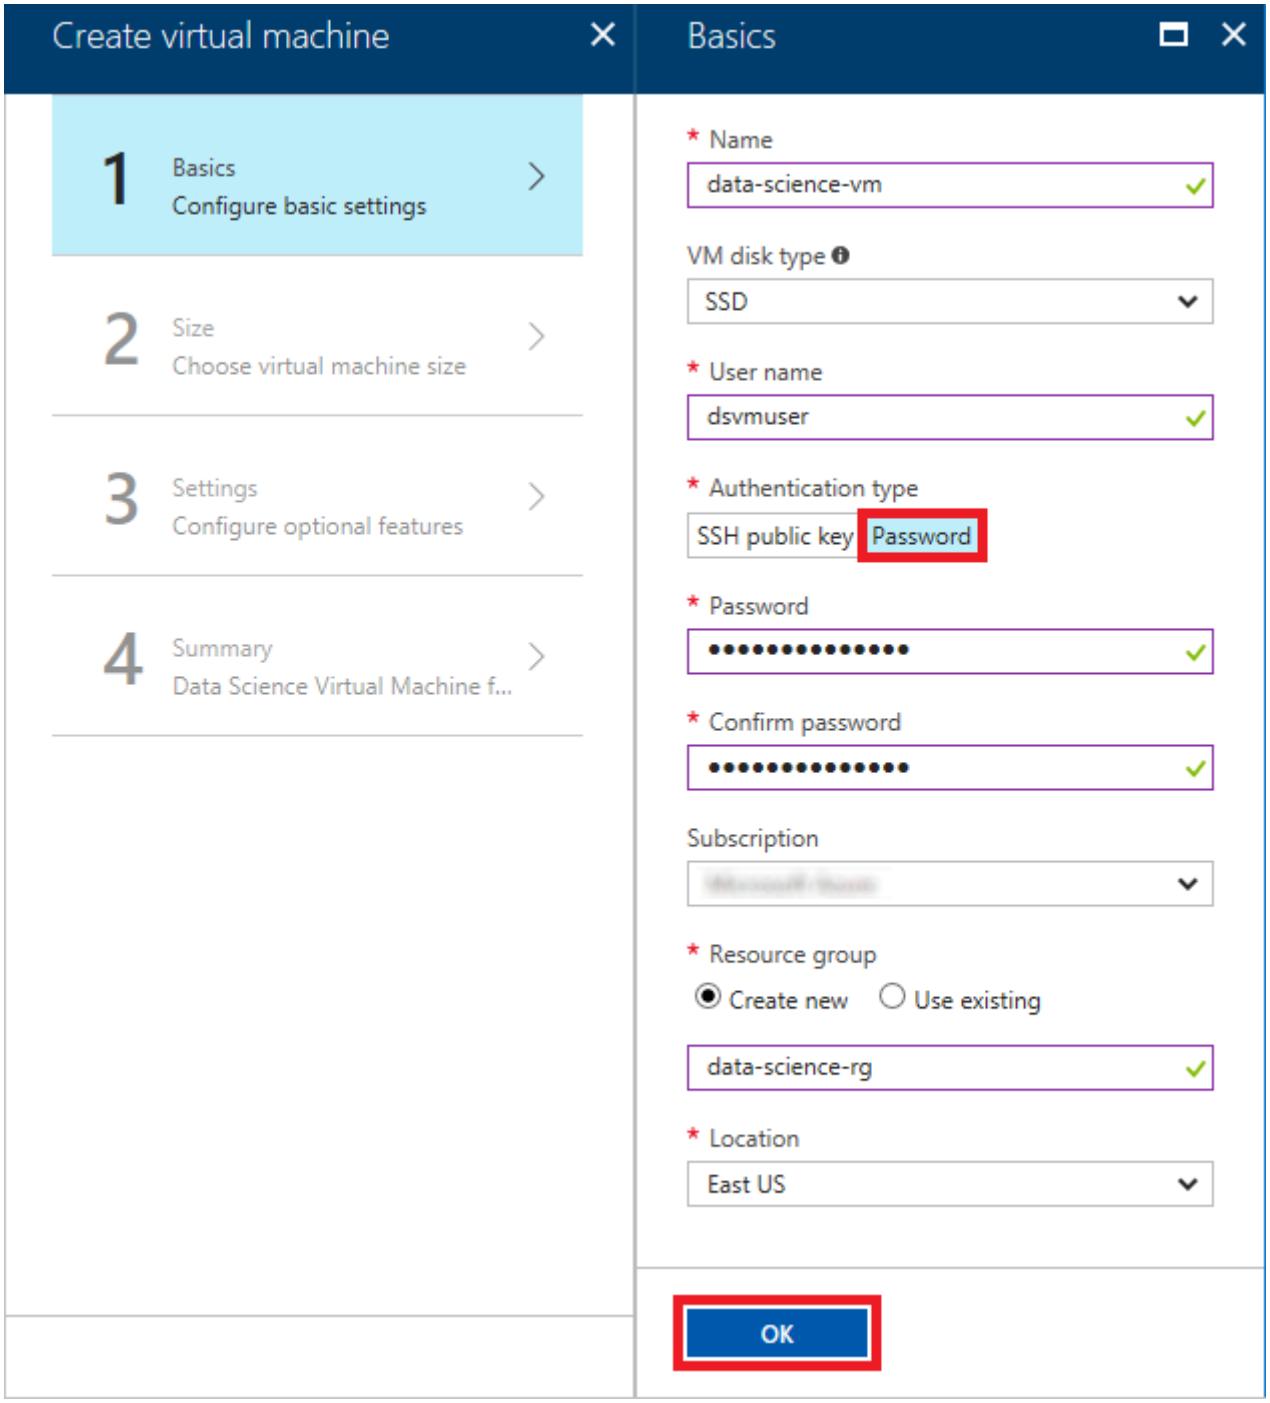

Entering basic information about the VM

5. Next, choose a size for the VM. In order to show all size options available, click **View All**. Scroll down and select **DS1_V2 Standard**, which provides a low-cost way to experiment with Data Science VMs. Then click the **Select** button.

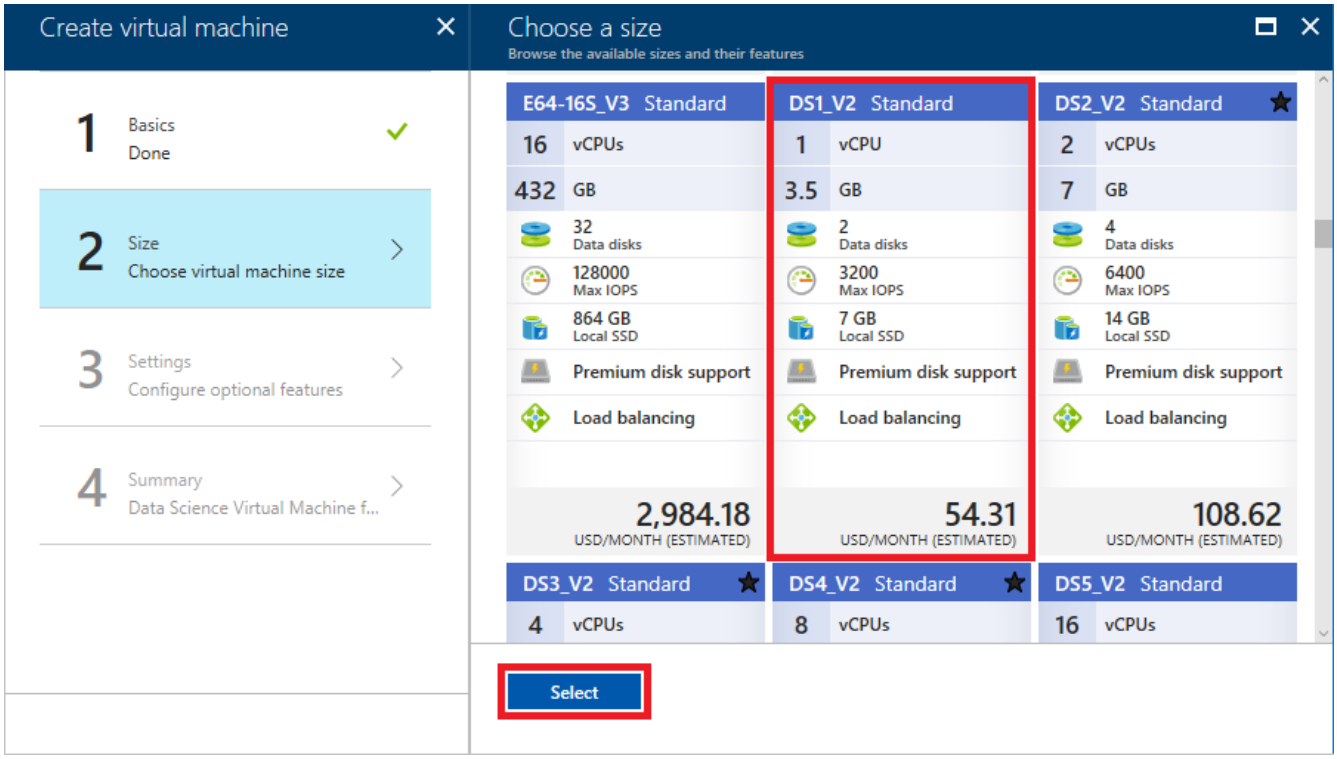

Choosing a VM size

6. Click **OK** at the bottom of the "Settings" blade. Then take a moment to review the options you selected for the VM, and click **Create** to create it.

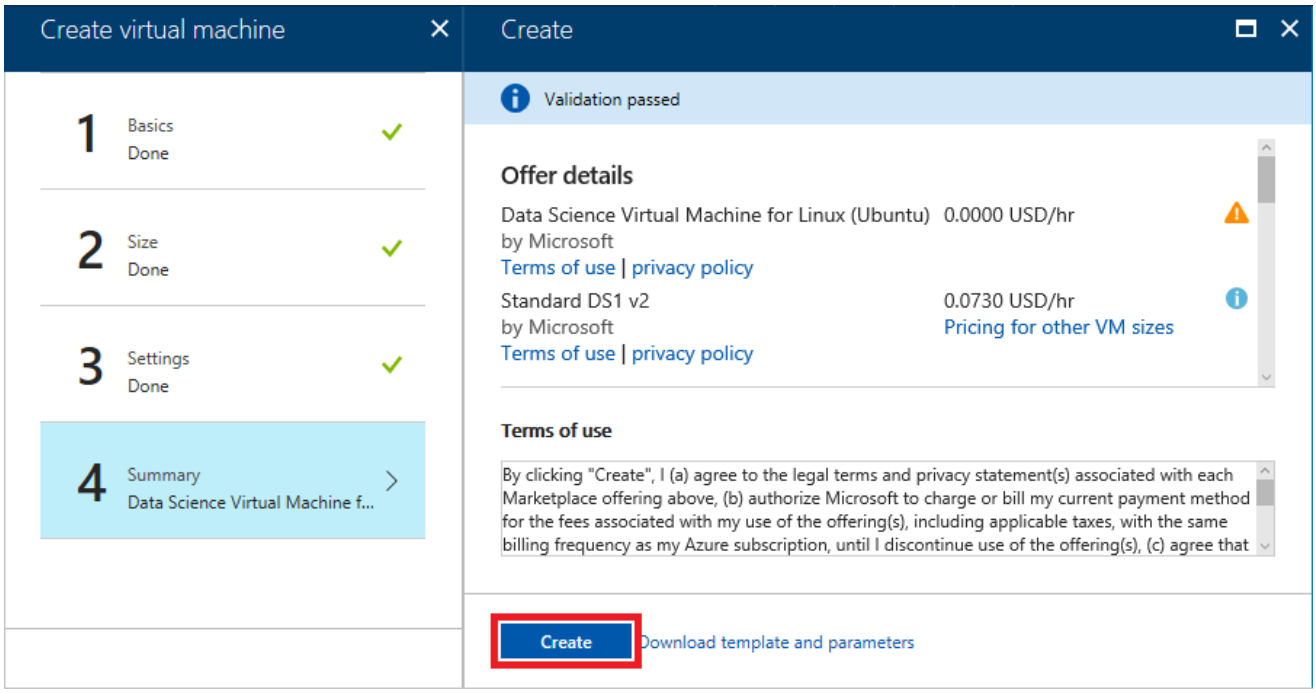

Creating the VM

Microsoft

7. Click **Resource groups** in the menu on the left side of the portal. Then click the resource group whose name you specified in Step 4.

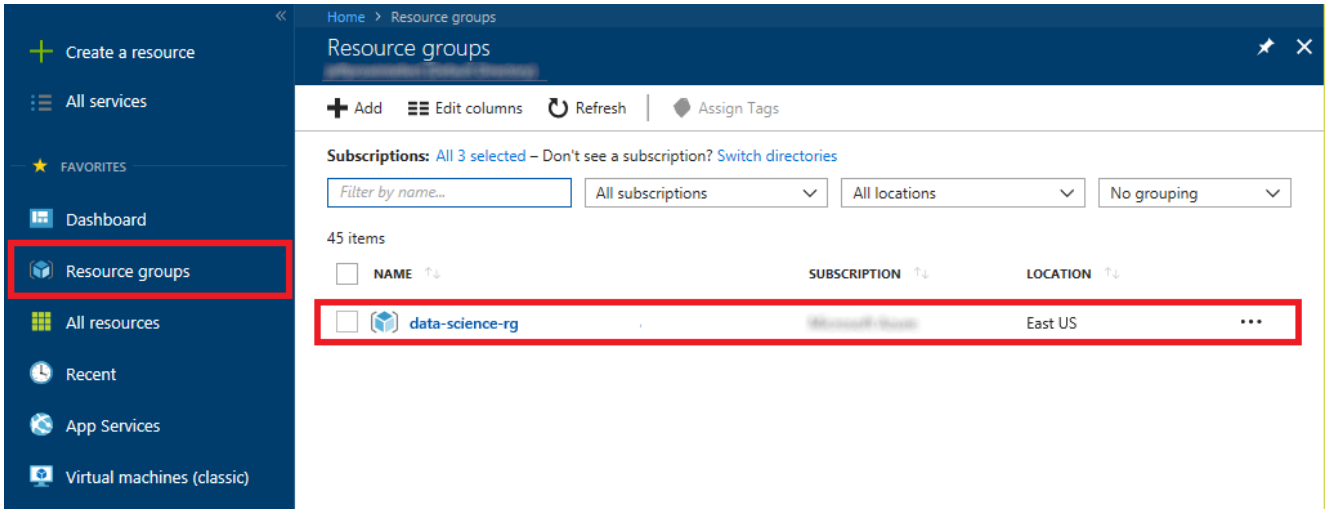

Opening the resource group

8. Wait until "Deploying" changes to "Succeeded" indicating that deployment has completed. Deployment typically takes 5 minutes or less. Periodically click **Refresh** at the top of the blade to refresh the deployment status.

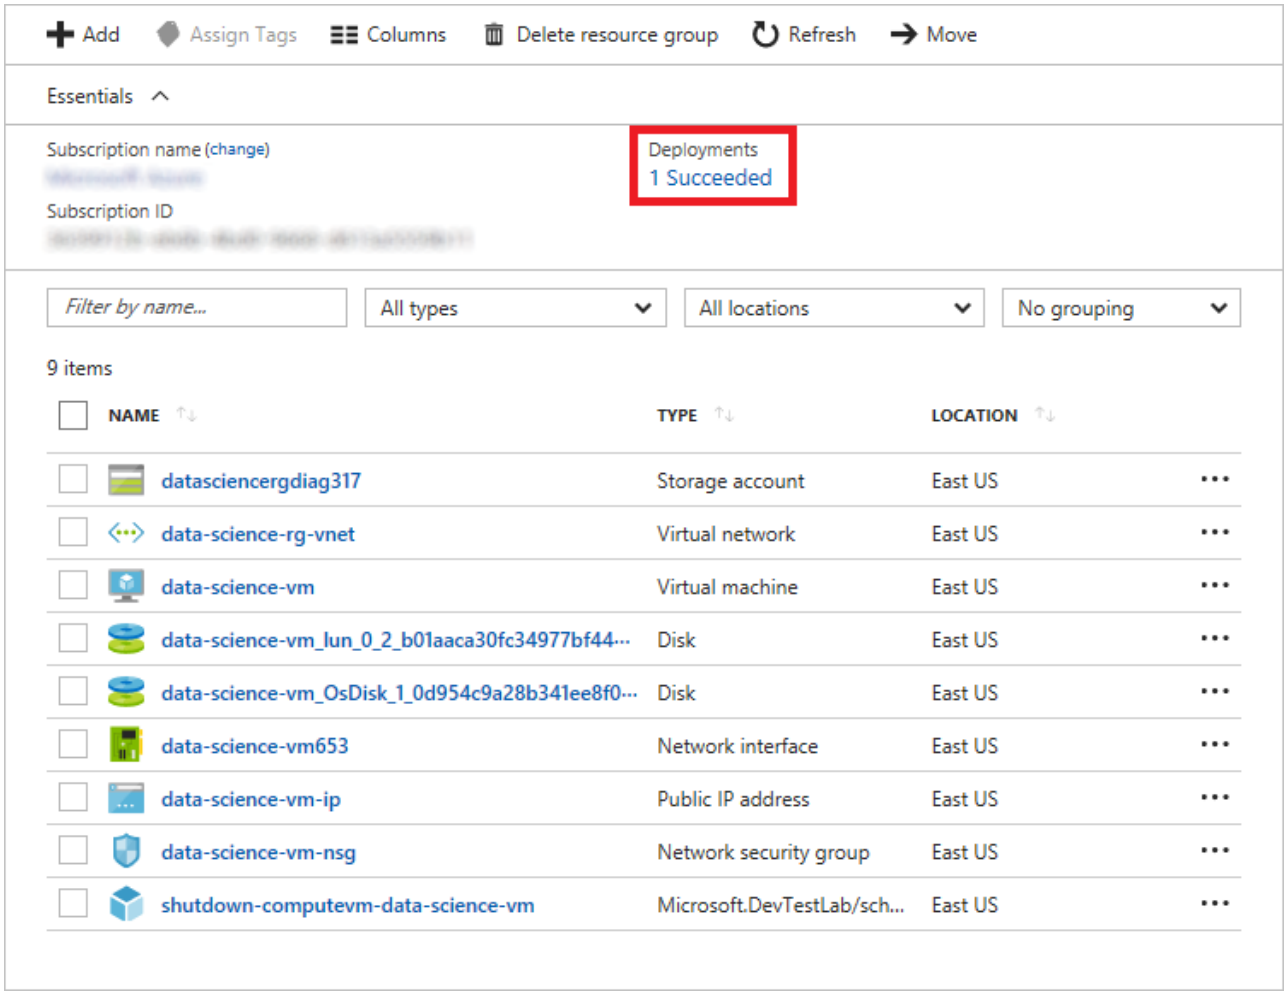

Monitoring the deployment status

The VM has been created. The next step is to connect to it remotely so you can work with the VM's Ubuntu desktop.

Exercise 2: Connect to the Data Science VM

In this exercise, you will connect remotely to the Ubuntu desktop in the VM that you created in the previous exercise. To do so, you need a client that supports Xfce, which is a lightweight desktop environment for Linux.

1. If you don't already have an Xfce client installed, download the X2Go client and install it now. X2Go is a free and open-source Xfce solution that works on a variety of operating systems. The instructions in this exercise assume you are using X2Go, but you can use any client as long as it supports Xfce.

2. Return to the Azure Portal and the blade for the resource group containing the Data Science VM. Then click the VM.

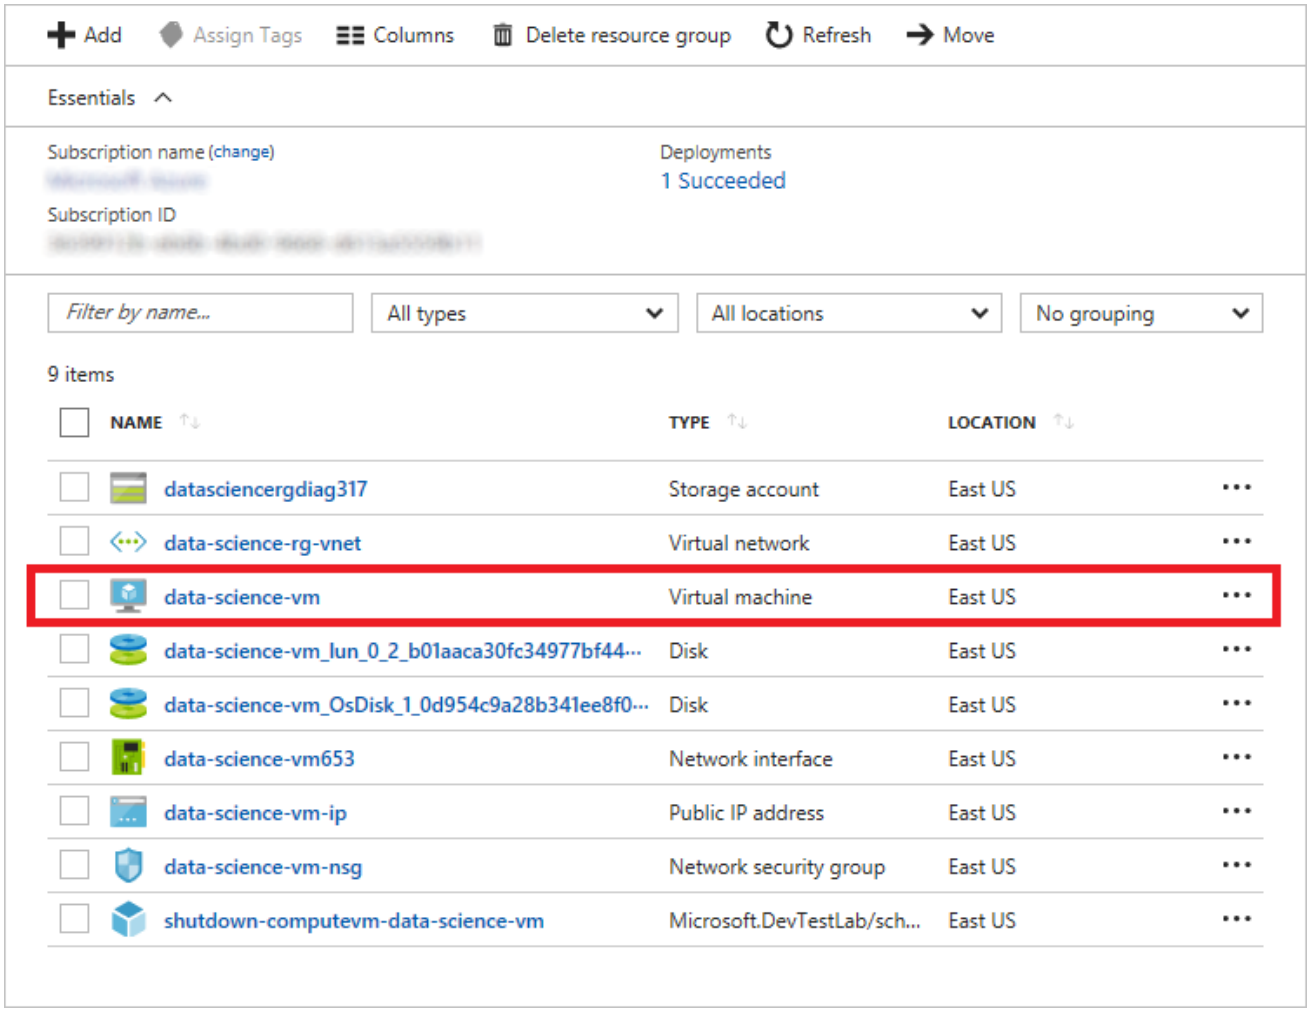

Opening the Data Science VM

3. Hover over the IP address shown for the VM and click the **Copy** button that appears to copy the IP address to the clipboard.

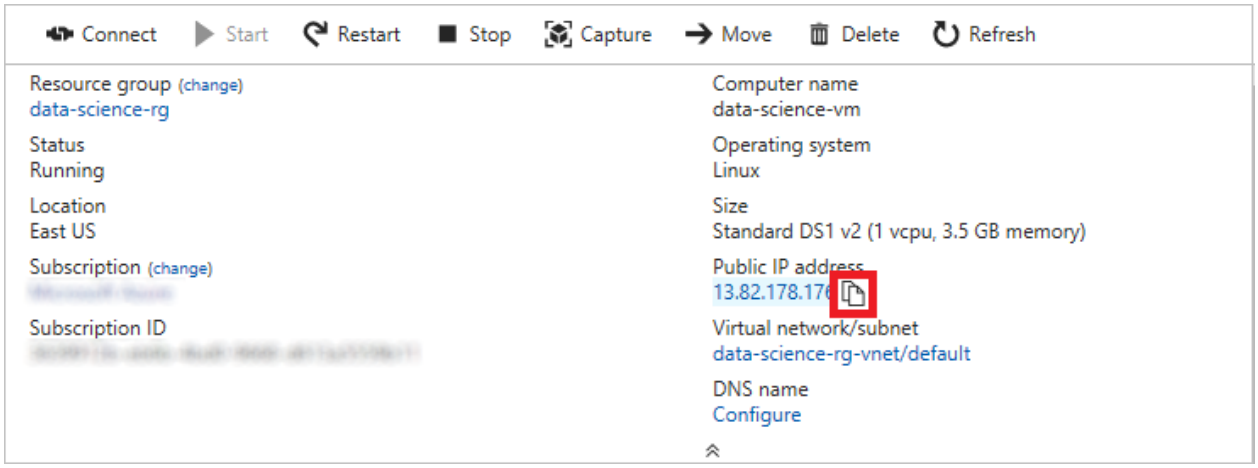

Copying the VM's IP address

4. Start the X2Go client and connect to the Data Science VM at the IP address that's on the clipboard using the user name you specified in the previous exercise. Connect via port **22** (the standard port used for SSH connections), and specify **XFCE** as the session type.

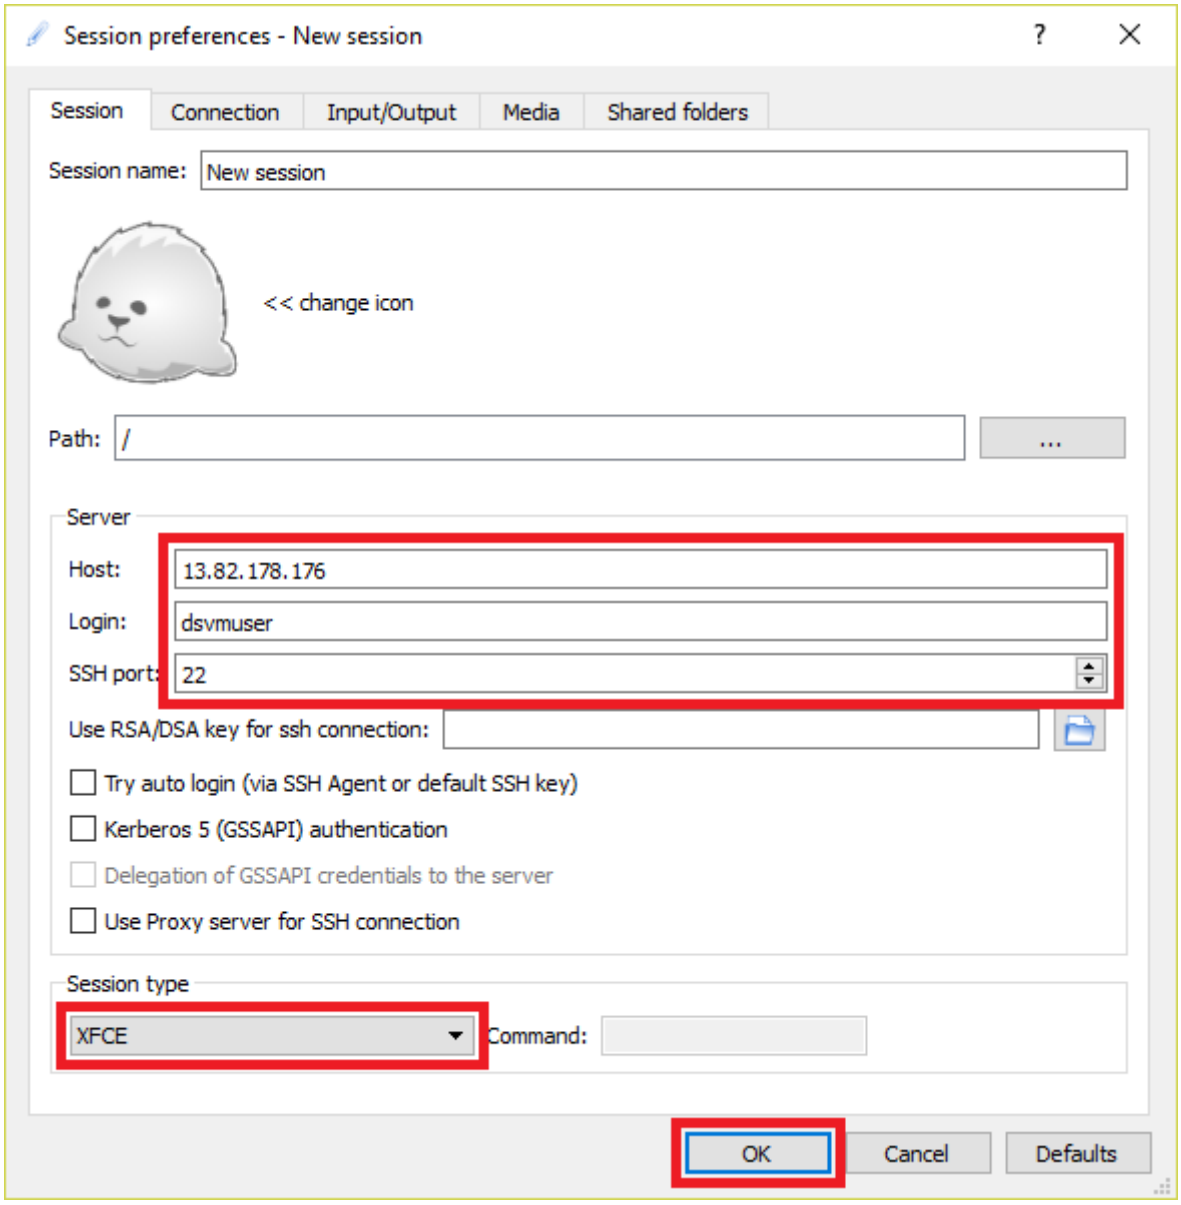

Connecting with X2Go

5. In the **New session** panel on the right, select the resolution that you wish to use for the remote desktop. Then click the **New session** panel.

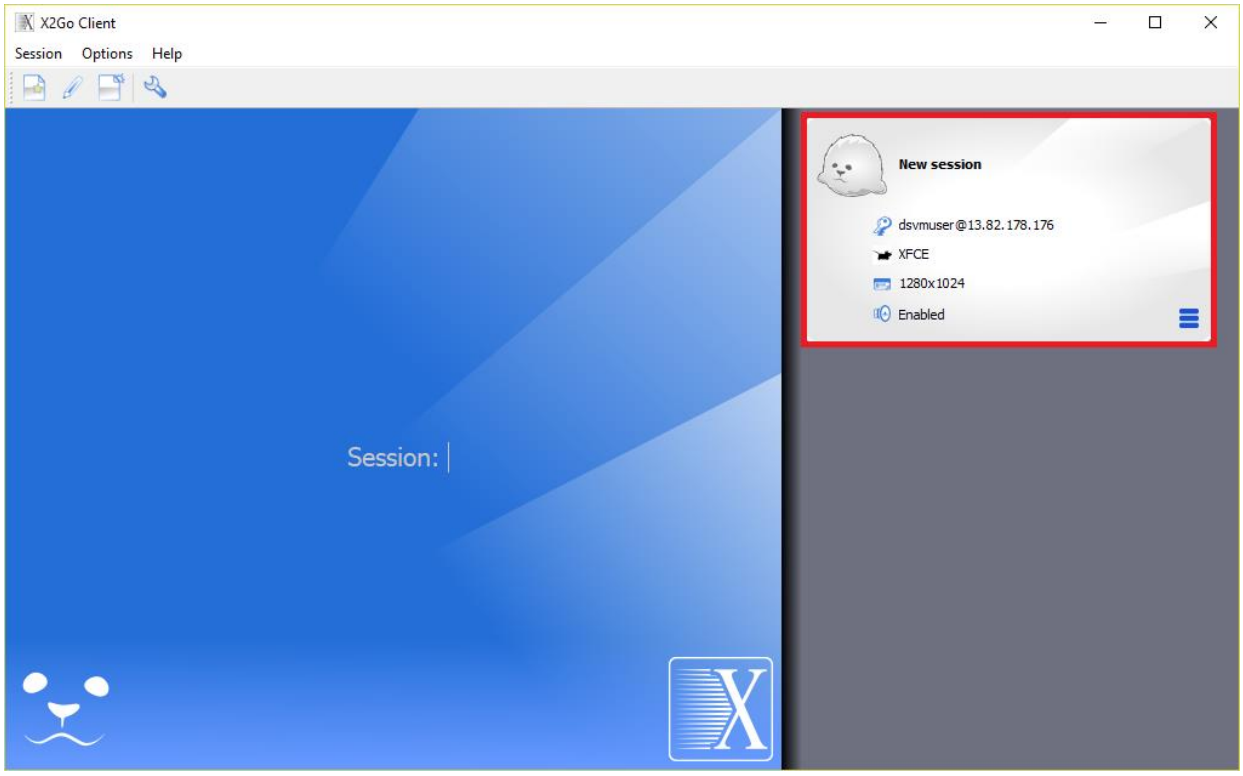

Starting a new session

6. Enter the password you specified in Exercise 1, and then click the **OK** button. If asked if you trust the host key, answer **Yes**. Also ignore any error messages saying the "SSH daemon could not be started."

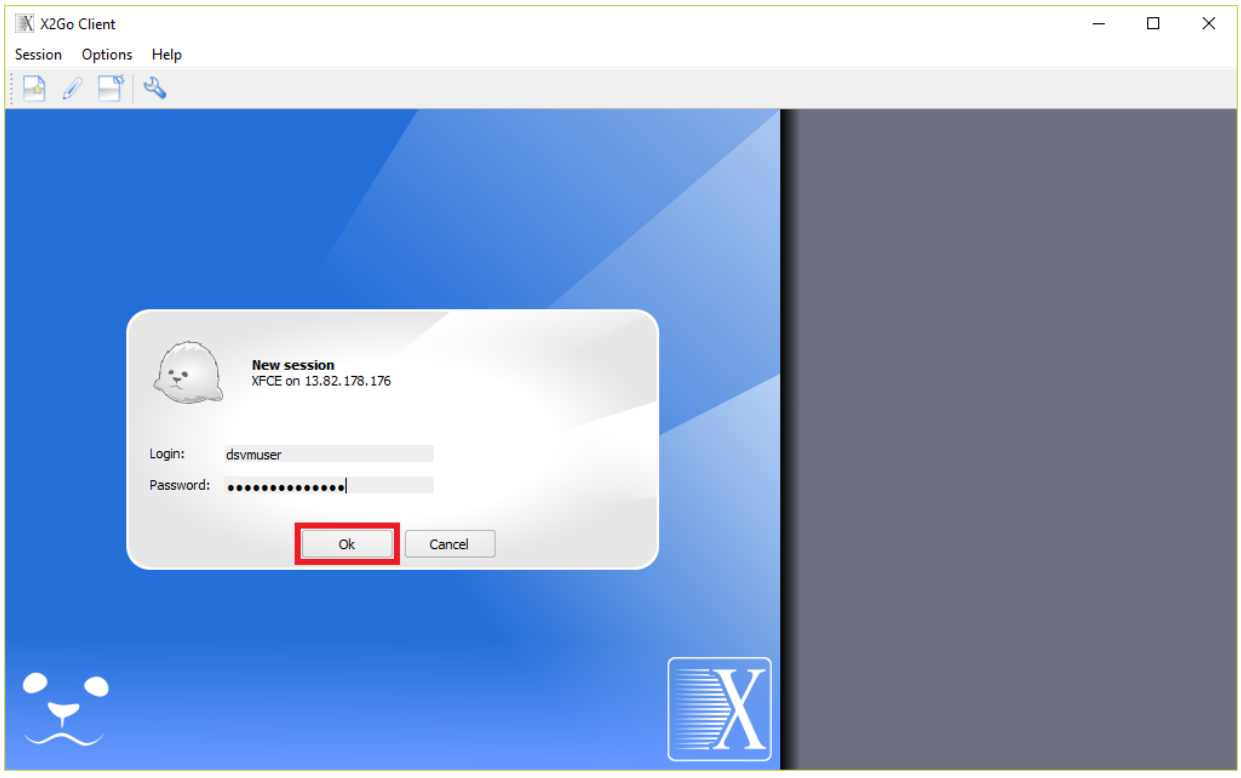

Logging into the VM

7. Wait for the remote desktop to appear and confirm that it resembles the one below.

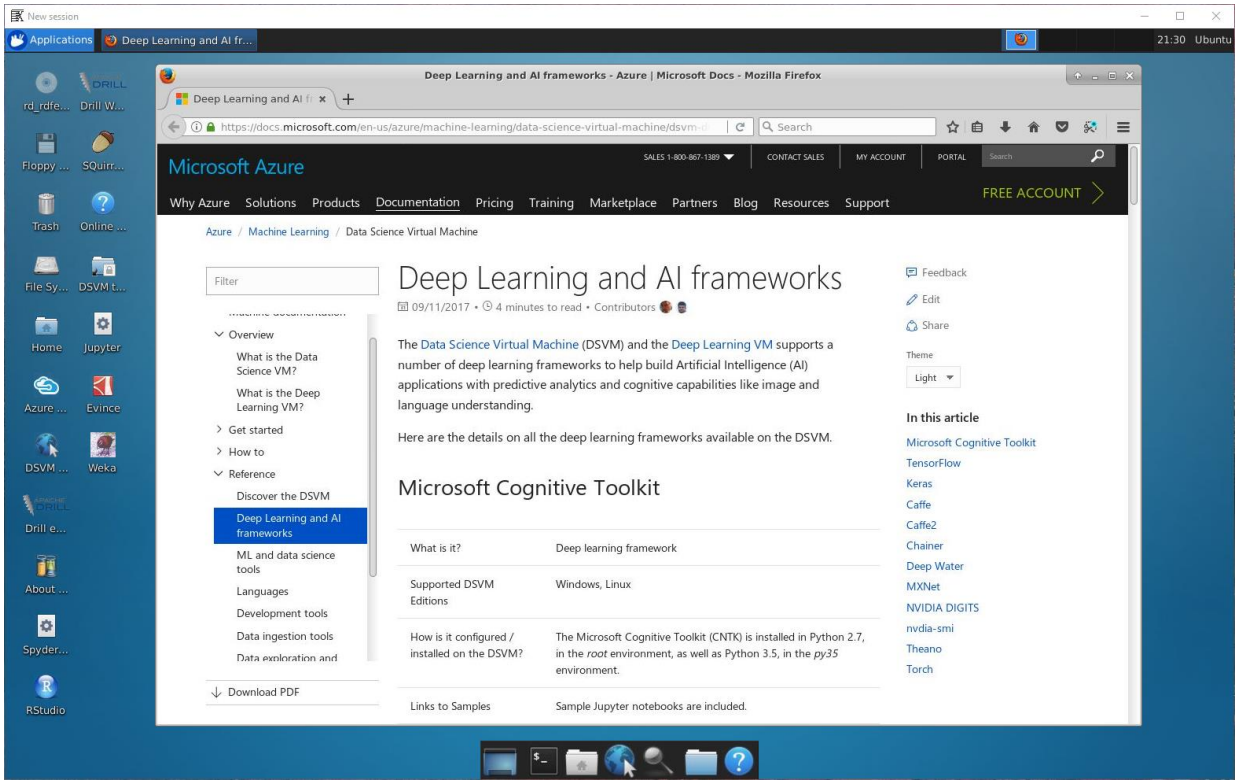

Connected!

Now that you are connected, take a moment to explore the shortcuts on the desktop. These are shortcuts to the numerous data-science tools preinstalled in the VM, which include Jupyter, R Studio, and the Microsoft Azure Storage Explorer, among others.

Exercise 3: Download a dataset and create a Jupyter notebook

In this exercise, you will import a dataset from Azure blob storage into the VM and load it into a Jupyter notebook. Jupyter is already installed in the VM and is accessible through the **Applications** menu or through the shortcut on the desktop. Jupyter notebooks are widely used in the data-science community to explore, transform, and visualize data. Notebooks are highly interactive, and since they can include executable code, they provide the perfect platform for manipulating data and building predictive models from it.

1. Click the Terminal icon at the bottom of the desktop to open a terminal window.

Opening a terminal window

2. Enter the following commands in the terminal window to create a "flights" subdirectory in the "notebooks" directory and download a dataset from Azure blob storage into the "flights" subdirectory:

```
cd notebooks
mkdir flights
cd flights
curl https://topcs.blob.core.windows.net/public/FlightData.csv --output flightdata.csv
```

6. Click **Applications** in the upper-left corner of the desktop. Then click **Development**, followed by **JupyterHub**.

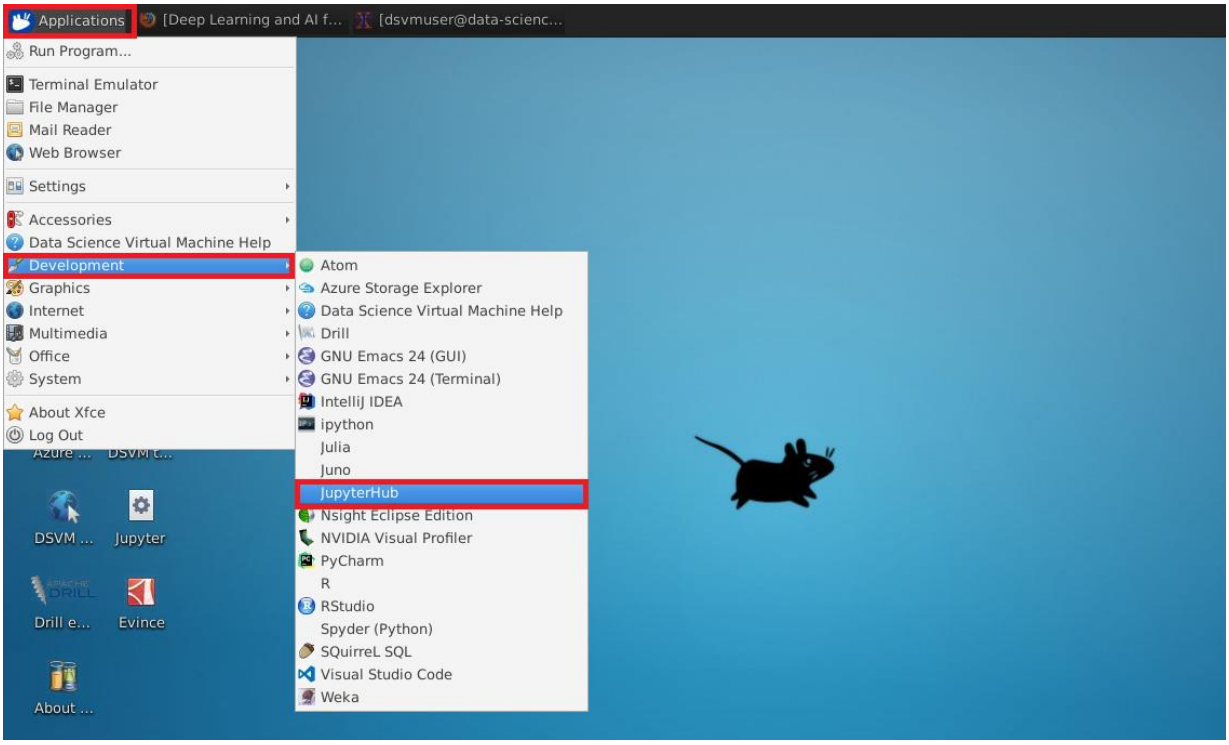

Launching JupyterHub

7. In the browser window that opens, click **flights** to open the "flights" directory.

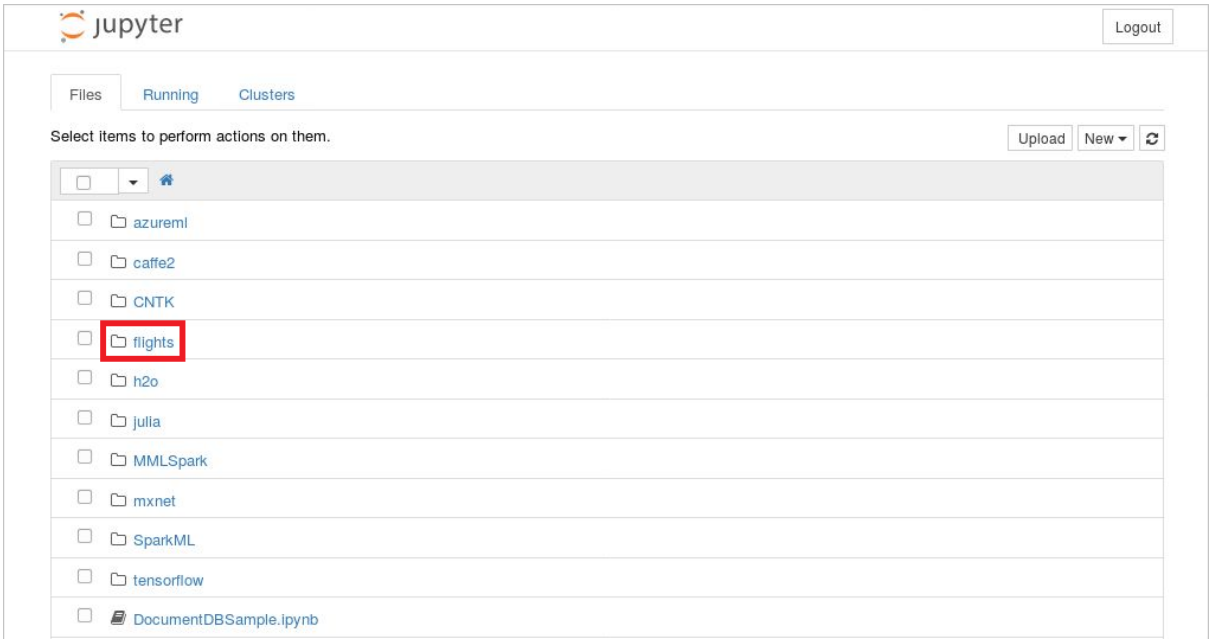

Opening the "flights" directory

8. Confirm that **flightdata.csv** is present in the "flights" directory. Then click the **New** button and choose **Python 3 Spark - local** from the drop-down list to create a new Jupyter notebook with a Python 3 kernel.

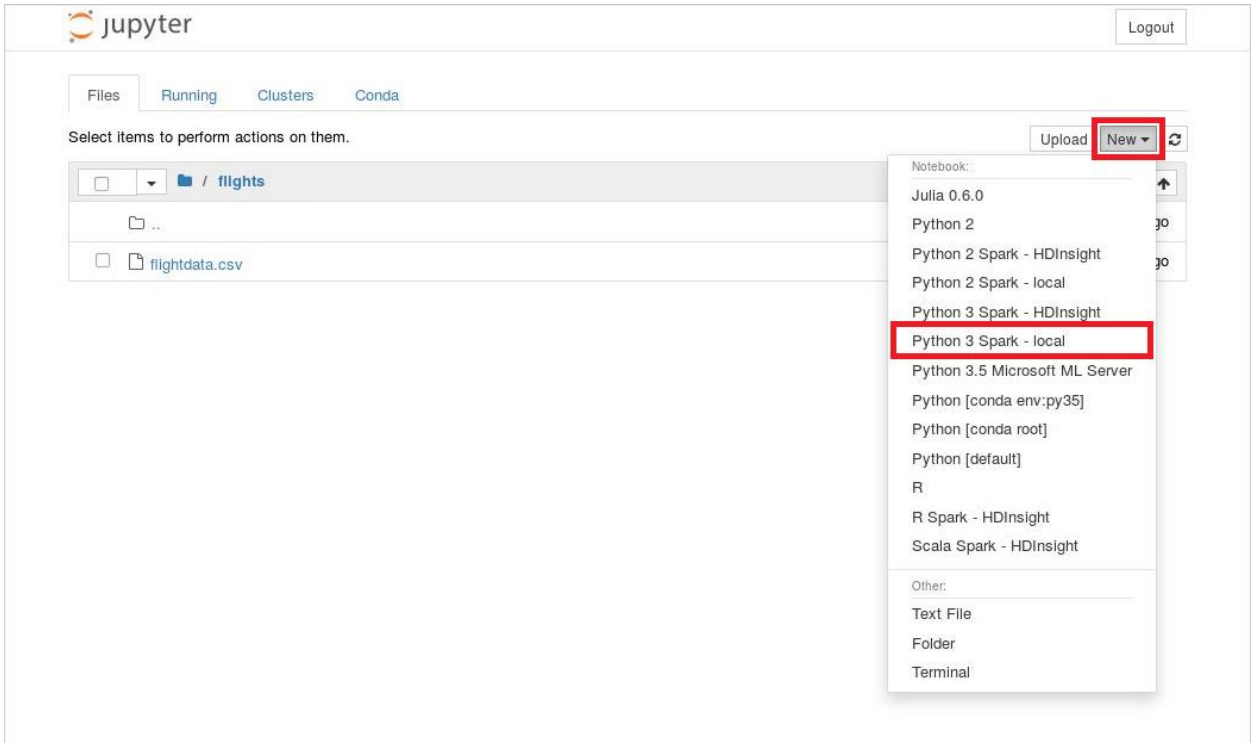

Creating a new Jupyter notebook

9. In the first cell of the notebook, enter the following Python code to load **flightdata.csv** and create a Pandas DataFramefrom it.

```
import pandas as pd

df = pd.read_csv('flightdata.csv')
df.head()
```

13. Select the **Run Cells** command from the **Cell** menu (or press **Ctrl+Enter**) to execute the Python code. Confirm that the output resembles the output below.

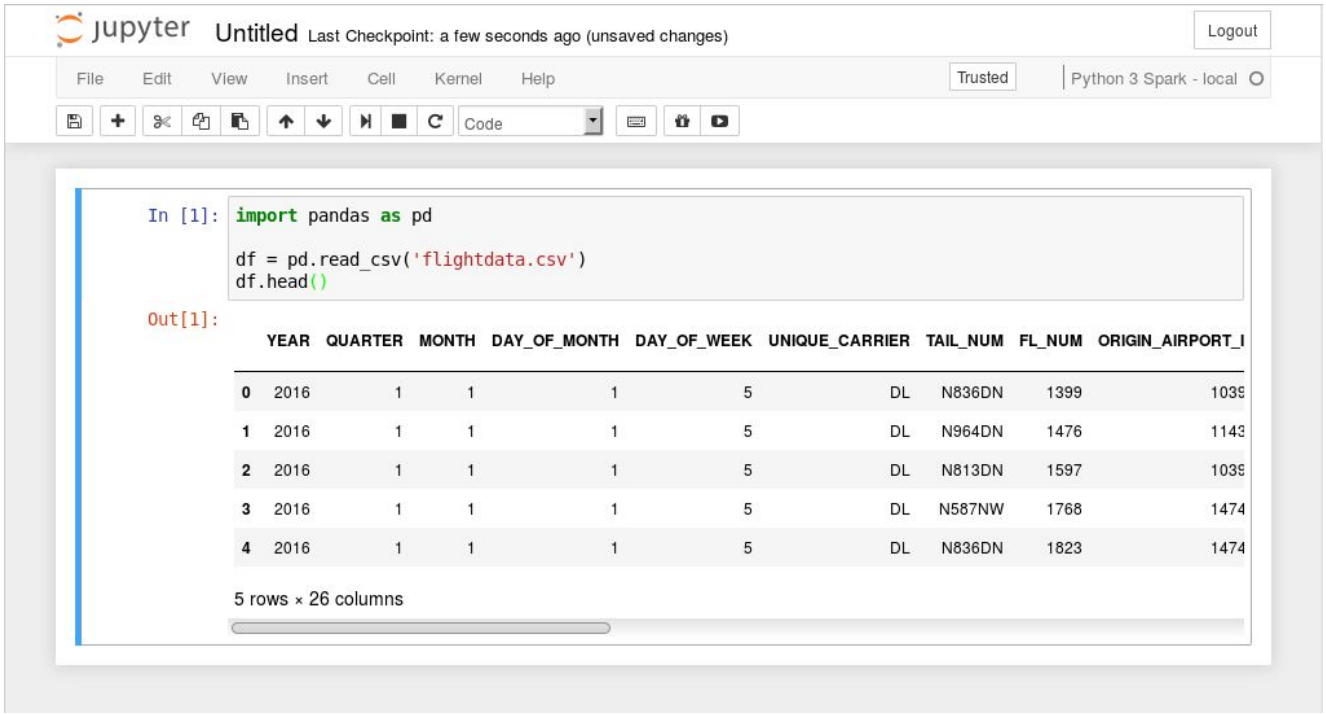

Loading the dataset

The DataFrame that you created contains on-time arrival information for a major U.S. airline. It has more than 11,000 rows and 26 columns. (The output says "5 rows" because DataFrame's head function only returns the first five rows.) Each row represents one flight and contains information such as the origin, the destination, the scheduled departure time, and whether the flight arrived on time or late. You will learn more about the data, including its content and structure, in the next lab.

14. Use the **File** -> **Save and Checkpoint** command to save the notebook.

15. Use the **File** -> **Rename...** command to name the notebook "FlightData."

If you check the "flights" directory, you should find that it now contains a file named **FlightData.ipynb** containing the Jupyter notebook you created. You will return to this notebook in the next lab and use it to prepare the data for use in a machine-learning model.